

UCLA BFB Alpha Research Competition

Sponsorship Proposal | Fall 2026 - Spring 2027

A research-first quantitative investing competition in which UCLA students turn an original market thesis into a documented, executable strategy and test the frozen submission on previously unseen data.

Competition at a Glance

- **Organizer:** Bruins in Finance and Banking at UCLA.
- **Format:** teams of up to three students complete research, historical testing, a strategy freeze, and an unseen forward-window evaluation.
- **Markets:** U.S. equities in a frozen S&P 500 universe and liquid Binance spot cryptocurrency markets.
- **Submission:** a research memo, reproducible code, dependency and data documentation, backtest results, and a frozen strategy configuration.
- **Evaluation:** research quality, backtest rigor, unseen forward-window performance, risk controls, and reproducibility.
- **Portfolio:** \$100,000 of simulated starting capital; no participant capital is put at risk.
- **Proposed schedule:** research from September through November 2026, followed by an extended forward window from December 2026 through early March 2027, ending before spring break; finalists and winners are presented in spring quarter.

Partnership Summary

Sponsorship begins at a minimum contribution of \$1,000. Terms are negotiated directly with each partner so support can match the partner's recruiting, research, education, or brand objectives without a public tier structure.

- Cash and in-kind support can fund prizes, market data, compute, research tooling, event costs, and student programming.
- Partners may contribute judges, technical mentors, data access, or a carefully scoped research prompt.
- Recognition and access are agreed in advance and can include event visibility, finalist presentations, and consent-based recruiting materials.
- BFB can coordinate an on-campus information session, research workshop, practitioner talk, or small-group recruiting event with the sponsor.

Competition Design

The competition is designed as a compact research pipeline rather than a short-horizon trading contest. Participants are rewarded for defensible reasoning, realistic testing, and disciplined implementation.

Participant workflow

- **1. Form a research question** — define the claimed source of edge, economic rationale, market, horizon, and expected failure modes.
- **2. Acquire and document data** — record provenance, timestamps, licensing constraints, transformations, and known limitations.
- **3. Build the signal and portfolio** — translate the thesis into deterministic entry, exit, sizing, exposure, and risk logic.
- **4. Run historical validation** — separate development from out-of-sample testing and include realistic trading frictions.
- **5. Submit and freeze** — deliver code, configuration, documentation, and declared dependencies before the deadline.
- **6. Complete the forward window** — run the unchanged strategy through the official evaluation environment and data feed.
- **7. Present and defend** — explain the result, attribution, limitations, and lessons to the judging panel.

Proposed timeline

Phase	Proposed timing	Primary output
Registration and orientation	September 2026	Team registration, rules briefing, and technical onboarding
Research and development	September-November 2026	Thesis, data work, implementation, and historical testing
Submission and freeze	Late November 2026	Research memo, code package, results, and locked configuration
Unseen forward window	December 2026-early March 2027	Official simulated execution through the period before spring break
Verification and finals	March-early Spring Quarter 2027	March reproduction checks; finalist presentations and awards in spring quarter

Required submission package

- Research memo stating the thesis, evidence, assumptions, portfolio construction, and limitations.
- Runnable strategy code with a pinned environment, configuration file, and clear entry point.
- Backtest report with returns, drawdowns, turnover, exposure, trade statistics, benchmarks, and sensitivity checks.
- Data dictionary and provenance record identifying the source, fields, timestamp convention, and preprocessing.
- Reproduction instructions and a signed confirmation that the post-freeze strategy will not be changed.

Eligible markets and data responsibility

- Equities use a frozen S&P 500 constituent universe specified by the organizers to reduce ambiguity and survivorship leakage.
- Crypto uses up to 1,000 Binance spot instruments selected under published liquidity and data-availability criteria.
- Participants are responsible for lawfully obtaining and disclosing research-phase data; official evaluation data and conventions are standardized by BFB.

Research and Historical Testing Standards

A strong submission must make it possible for a judge to understand what is being predicted, why the effect could exist, how it becomes a portfolio, and whether the reported evidence survives reasonable attempts to break it.

Research memo

- **Thesis and mechanism:** state the hypothesized source of return and why it should persist after costs.
- **Data provenance:** identify each source, field, timestamp, publication lag, transformation, and licensing restriction.
- **Signal construction:** define the feature, forecast, rebalance frequency, lookback windows, and treatment of missing or stale observations.
- **Portfolio construction:** explain security selection, sizing, diversification, cash treatment, leverage, and rebalancing.
- **Execution assumptions:** document order timing, prices used, fees, spread or slippage, liquidity filters, and market-hours conventions.
- **Validation plan:** distinguish model development, parameter selection, validation, and untouched out-of-sample evidence.
- **Risk and limitations:** identify concentration, crowding, regime, liquidity, capacity, model, and operational risks.

Historical testing

- Use chronological train, validation, and out-of-sample segments; random shuffling is generally unsuitable for time-series evidence.
- Use only information that would have been available at each decision timestamp, including publication and constituent dates.
- Address survivorship bias, delistings, splits, dividends, symbol changes, stablecoin or venue changes, and other relevant market events.
- Apply transaction costs and slippage that match instrument liquidity and strategy turnover; disclose any zero-cost result separately.
- Report parameter sensitivity and adjacent specifications so a narrow optimum is visible rather than hidden.
- Show subperiod, regime, and instrument-level attribution; aggregate performance alone is not sufficient.
- Compare with a relevant benchmark and simple baselines, and separate market beta from the claimed alpha source.
- Report return, volatility, drawdown, turnover, exposure, hit rate, trade count, and the distribution of outcomes—not Sharpe alone.

Data and reproducibility

- The final package must run from documented inputs without undisclosed manual adjustments or post-freeze overrides.
- Random seeds, package versions, configuration values, and data joins must be fixed and recorded.
- Judges may request an execution trace, intermediate signal sample, or simplified reproduction case during finalist verification.

Evaluation Framework

Statistical interpretation. The extended December-to-March forward window provides more evidence than a five-week test, but it remains too short to make an annualized Sharpe or Calmar ratio statistically decisive, especially for low-turnover equity strategies. Forward performance therefore complements the broader research and historical record.

Proposed scoring

Component	Weight	What judges assess
Research thesis and implementation	30%	Economic rationale, originality, signal clarity, and faithful translation into portfolio rules
Backtest rigor and robustness	35%	Leakage control, realistic frictions, out-of-sample design, sensitivity, attribution, and benchmark quality
Unseen forward-window performance	20%	Absolute and risk-adjusted result, consistency, drawdown, execution quality, and behavior versus expectation
Risk controls and reproducibility	15%	Compliance, concentration and exposure control, code quality, documentation, and successful reproduction

How the forward window is interpreted

- No single annualized ratio or approximately three-month return determines the winner.
- Judges compare realized behavior with the strategy's stated thesis, expected turnover, exposure, and risk budget.
- Drawdown path, concentration, implementation errors, and unexplained deviations can outweigh a strong headline return.
- Trade count and independence are considered before interpreting Sharpe, Calmar, hit rate, or similar metrics.
- Low-frequency strategies may remain eligible for research awards even when the forward sample is too small for a performance award.

Judging and verification

- Judges score submissions independently against a written rubric before panel discussion.
- The composite score establishes the initial ranking; the panel may adjust only for documented rule, integrity, or reproduction findings.
- Finalists may be asked to rerun a designated period, explain individual trades, and reconcile reported figures with official logs.
- Conflicts of interest are disclosed, and an affected judge does not score the relevant team.
- BFB publishes the final rubric, data conventions, and tie-break procedure before the strategy freeze.

Proposed recognition

- Overall: first, second, and third place.
- Research awards: Best Research Thesis; Best Backtest and Validation; Best Risk Management.
- Performance awards: Best Forward-Window Performance; Best Risk-Adjusted Strategy.
- Partner recognition: Sponsor's Choice, based on a disclosed sponsor brief or finalist presentation.

Portfolio, Execution, and Risk Controls

The baseline rules below make results more comparable and reduce the chance that a short sample rewards a single concentrated bet. They remain proposed rules until the final handbook is published before launch.

Proposed hard controls

Control	Proposed baseline	Application
Single-position size	Maximum 20% of portfolio value	Measured after each rebalance; excess exposure must be reduced at the next permitted execution
Gross exposure	Maximum 150%	Long absolute exposure plus short absolute exposure; derivatives, if allowed, use defined notional treatment
Net exposure	-100% to +100%	Measured consistently across instruments and cash; no hidden leverage through offsetting positions
Forward drawdown	Maximum 35%	Breach removes eligibility for overall and performance-based awards, subject to incident review
Risk-adjusted award sample	At least 30 independent completed trades, including at least 10 in the forward window	Applies to metric-based awards; trades split mechanically to inflate the count are consolidated

Execution model

- All portfolios begin with the same simulated capital and use published fees, spread or slippage, and fill conventions.
- Orders are timestamped and evaluated against data available at the decision time; same-bar look-ahead is prohibited.
- The organizer may apply liquidity screens, participation limits, or rejected-fill logic consistently across teams.
- Corporate actions, delistings, halted instruments, symbol changes, and missing prices are handled under a written market-specific policy.
- Any organizer-side outage or bad-data event is logged and reviewed before a result is included in the official record.

Performance-award eligibility

- There is no universal minimum trade count to enter the competition or qualify for research awards.
- Metric-based awards require a sufficient and economically independent sample; repeated partial fills do not create new trades.
- A risk-limit breach, prohibited data use, material post-freeze change, or failed reproduction may remove performance eligibility.
- The judging panel retains discretion for exceptional market or platform events, but any exception is documented and applied consistently.

Specialized Research Designation

- A team using a single asset, infrequent event strategy, unusually long holding period, or unique data source may request the designation before the freeze.
- The request must explain why the general concentration or trade-count rule would distort the research design and propose an alternative risk limit.
- Approved teams remain eligible for thesis, validation, and risk-management awards; metric-based awards may be unavailable when the sample is too small.
- Designation decisions are recorded before evaluation so exceptions cannot be granted after performance is known.

Technical Operations and Governance

A credible competition depends on an auditable chain from data and code to orders, positions, and final metrics. The operating model therefore prioritizes reproducibility, consistent treatment, and clear incident handling.

Data and environment

- BFB publishes eligible universes, bar frequency, timezone, calendars, price fields, corporate-action treatment, and benchmark definitions.
- Official evaluation uses organizer-controlled market data and a fixed execution interface; research-phase external data must be disclosed.
- Dependencies are pinned in an approved environment. Network access and undeclared external calls can be disabled during official runs.
- Teams receive a validator or sample dataset before the deadline so schema, output, and runtime errors can be corrected before freezing.

Submission, freeze, and logging

- Each submission receives a timestamped archive and cryptographic hash covering source code, configuration, dependencies, and documentation.
- The frozen package is the only package used for the forward window; emergency fixes require organizer approval and are fully disclosed.
- Official runs record inputs, signals, orders, fills, positions, cash, exposure, exceptions, and calculated performance metrics.
- Automated checks flag missing data, stale signals, exposure breaches, non-deterministic output, runtime failure, and inconsistent accounting.

Incident handling

- Organizer-side data or platform incidents are separated from participant code failures and assessed under a written correction policy.
- If a period is rerun, the same rule is applied to all affected teams and both the original and corrected records are retained.
- Material rule interpretations are communicated to all participants, with an opportunity to raise questions before they affect scoring.

Finalist verification

- BFB reproduces headline metrics from the archived package and official logs.
- Finalists reconcile return, P&L, drawdown, exposure, and trade count to a common calculation methodology.
- Judges may inspect a sample of signals and trades for timestamp integrity, rule compliance, and consistency with the research memo.
- Plagiarism, fabricated results, unlicensed data, undisclosed manual intervention, or deliberate leakage can result in disqualification.

Intellectual property and sharing

- Participants retain ownership of their original research and code, subject to the competition rules and any third-party data terms.
- Sponsors receive only the materials and recruiting information described in the agreed partnership scope.
- Resume books, contact details, source code, or non-public research are shared only with participant consent and appropriate permissions.

Sponsor Participation

Partnerships are intentionally flexible. BFB works with each sponsor to define a useful combination of competition support, practitioner involvement, campus engagement, and recognition while preserving independent judging and participant consent.

Contribution policy

- Minimum sponsorship is \$1,000; all terms, deliverables, and recognition are negotiable.
- Support may be financial or in kind, including data licenses, cloud credits, software access, venue support, or practitioner time.
- BFB does not publish fixed sponsorship tiers. A short written agreement records the contribution, intended use, deliverables, timing, and contacts.

Ways a partner can contribute

- Prize and program funding that expands participation and recognizes high-quality research.
- Market data, alternative data, compute, research tooling, or educational access for approved competition use.
- Judges or mentors with experience in systematic investing, data science, execution, portfolio construction, or risk.
- A scoped research prompt, dataset demonstration, or technical workshop that remains optional and does not advantage one team.
- Recruiting and campus programming coordinated with BFB and consistent with UCLA policies.

Potential sponsor involvement and deliverables

- Recognition in agreed competition materials, event slides, and finalist programming.
- Opportunity to nominate a judge, attend finalist presentations, and present a Sponsor's Choice recognition.
- Access to an opt-in resume book or follow-up introductions where participants have provided consent.
- A post-competition summary covering participation, research themes, event delivery, and agreed program outcomes.
- Research code or confidential strategy detail is not transferred unless a participant separately agrees.

On-campus engagement

- Information session: introduce the firm, investment approach, teams, internships, and recruiting process.
- Research workshop: walk through idea formation, validation, portfolio construction, or common backtesting failures.
- Practitioner panel or fireside chat: discuss careers and current research questions with a mixed student audience.
- Small-group format: office hours, coffee chats, or finalist feedback sessions for deeper technical interaction.
- Recommended format: a 45-minute practitioner talk, 15-minute competition briefing, and 30-minute networking session.

Use of sponsorship

- Participant prizes and recognition; data and compute access; technical infrastructure; venue and event operations; educational programming.
- BFB can agree a designated use in advance and provide a concise post-event account of delivery.

Working process

- **1. Introductory conversation** — confirm the sponsor's objectives, timing, contribution type, and relevant internal stakeholders.
- **2. Scope the partnership** — agree involvement, campus programming, recognition, participant access, and any data or compliance constraints.
- **3. Document and deliver** — record the arrangement, assign owners, and coordinate assets, speakers, judges, and event logistics.
- **4. Close the loop** — share agreed outcomes and discuss whether to continue the relationship in a future competition cycle.

Next step: BFB would welcome a short conversation to tailor a partnership around the sponsor's research, recruiting, or campus-engagement priorities.